

Curriculum Vitae Daniel David Kovacs

Software Architect | Lead Data Engineer
Winterthur, Switzerland (B permit)

Born: 07 April 1981, Budapest, Hungary
Email: daniel.david.kovacs@gmail.com
Phone: +41 79 309 26 21
GitHub: github.com/kicsikrumppli
LinkedIn: linkedin.com/in/danieldavidkovacs



Professional Summary

Hands-on software architect, and lead data engineer. I have over 11 years of experience in designing and implementing scalable software and data solutions. I have a keen interest in machine learning. I am skilled in software engineering, data modelling, and cloud platforms. Experienced in designing efficient data pipelines, messaging architectures, managing complex projects across industries including insurance, finance, life sciences, and transportation.

Skills

Data Engineering: Python, PySpark, Databricks, Kafka, Airflow, Polars, Pandas
Cloud Technologies: Azure (Databricks, Function Apps, CosmosDB, Azure ML), AWS (EMR, S3, Lambda, EC2, EKS)
Software Development: Python, Java, Scala
Tools, Frameworks, Databases: FastAPI, Spring Boot, PostgreSQL, CosmosDB, Neo4J, Docker
Languages: Hungarian (Native), English (C2), German (B2+)

Professional Experience

EPAM, Switzerland		Apr 2022 – Present		Software Architect & Lead Data Engineer
EPAM, Hungary		May 2014 – Apr 2022		Software Architect & Lead Data Engineer

Lead Data Engineer for Global Reinsurance Leader

Implemented ETL pipelines in a medallion architecture for financial data on Azure Databricks.

- Designed transformation-by-configuration library, enabling scaling to multiple use cases
- Improved pipeline resiliency and maintainability through modular design
- Established project onboarding process

Technologies: Azure Databricks, PySpark, Python, Azure Data Factory, Azure SQL

Solutions Architect for Global Reinsurance Leader

Designed and led the implementation of a cloud-native multi-tenant self-service data ingestion platform for insurance contract management with AI-driven intent and attribute extraction, and end-to-end lifecycle governance.

- Authored architecture documentation, securing project approval and successful delivery
- Designed, choreographed event-driven pipelines for on-demand, self-service ingestion across multiple business units
- Implemented append-only event sourced data model for replayable workflows and flexible downstream integrations
- Enforced governance with automated and manual validations, full audit trails and fine-grained RBAC

Technologies: Azure Event Grid, CosmosDB, Function Apps, Key Vault, Entra ID, Spring Boot

Lead Data Engineer for a Major Swiss Insurer

Implemented cloud migration and general optimizations of existing in-house data transformation pipelines used by the actuarial team for life insurance risk management.

- Migrated local Python pipelines to Azure ML, reducing model evaluation times by 90%
- Improved pipeline scalability via Pandas-to-Polars migration
- Enhanced DAG visualisation backend, accelerating insurance model development cycle

Technologies: Python, Numba, Pandas, Polars, Azure ML, Azure Blob Storage, Docker

Solutions Architect for a Business Information and Media Company

Scaled proof of concept into a production-ready commercial product of on-premises industrial computer vision solution to monitor mining conveyor belts for foreign object alerting, and production reporting. Designed event driven architecture with modular image processing pipeline.

- Implemented modular image processing pipeline on a ZeroMQ backbone separating video sampling, preprocessing, monitoring, and image segmentation and recognition phases; achieved a 10x performance increase in processed frames per second
- Designed resilient messaging for object detection event handling on RabbitMQ
- Integrated MLflow for end-to-end model training, versioning, deployment and inference serving
- Authored solution architecture document, which turned around high risk of delivery into a positive assessment

Technologies: Python, RabbitMQ, ZeroMQ, OpenCV, SQLAlchemy, MLflow

Lead Data Engineer for International Pharmaceutical R&D Company

Implemented components of a self-service data platform for drug research and discovery, providing cataloging, transformation and lineage of ingested biomedical database corpuses.

- Created Python DSL for schema and lineage management, shortening development cycles
- Enhanced Airflow execution, cutting down ingestion runtime costs with async processing
- Implemented systematic data quality checks

Technologies: Python, Airflow, PostgreSQL, Neo4J

Education

- BSc in Software Engineering – Budapest University of Technology, 2013
- MSc in Architecture and Engineering – Budapest University of Technology, 2007
- Master in Solar Energy Engineering – Högskolan Dalarna, Sweden, 2004